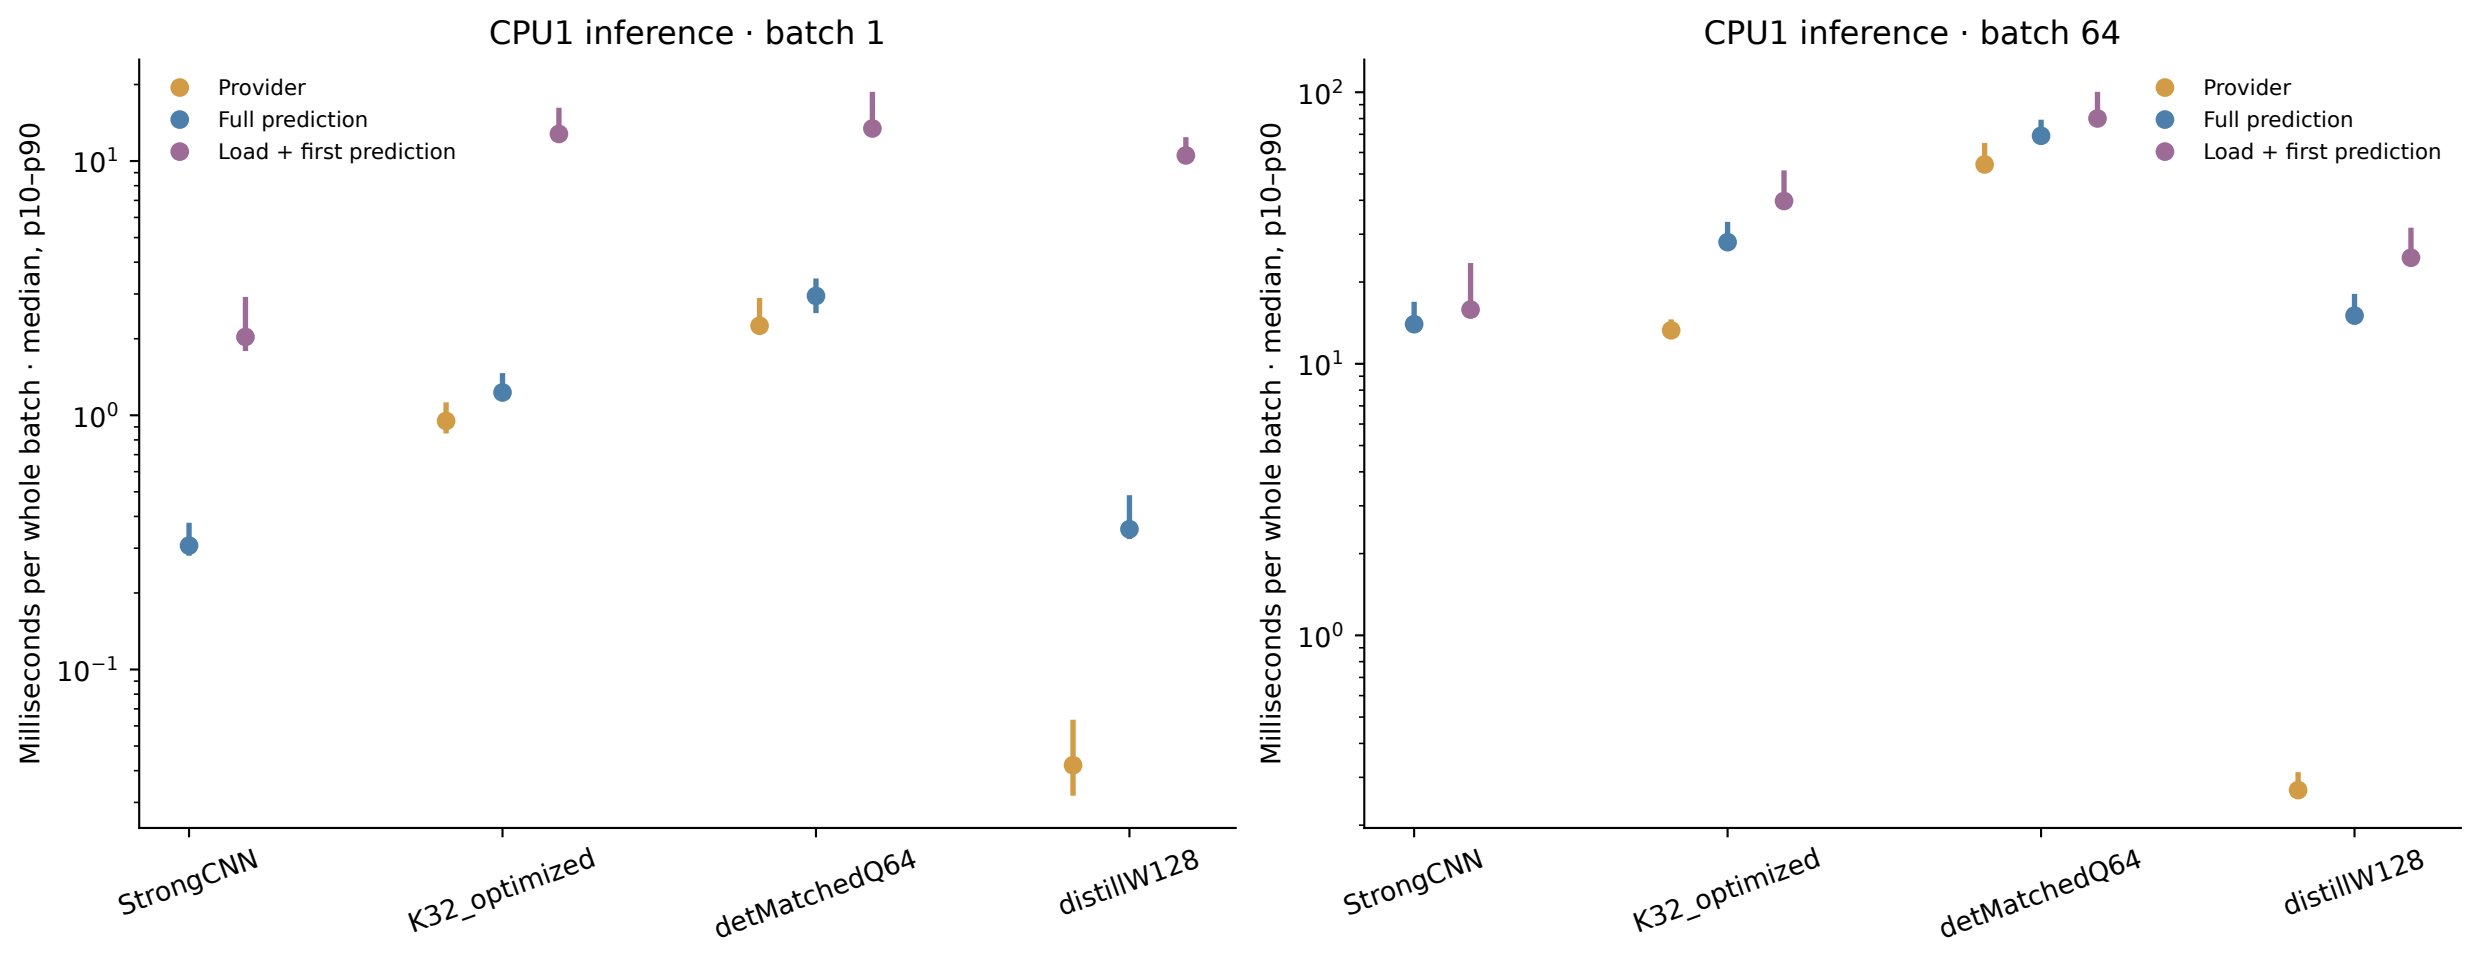
CPU1 inference · batch 1
CPU1 inference · batch 64
Provider
Full prediction
Load + first prediction
Milliseconds per whole batch · median, p10–p90
10^{2}
10^{1}
10^{0}
10^{-1}
StrongCNN
K32_optimized
detMatchedQ64
distillW128